

# Homework 1

Nicholas Allen

## Table of contents

Question 1	2
Question 2	3
Question 3	4
Tip	6

**Appendix** **10**

[Link to the Github repository](#)

---

**!** Due: Fri, Jan 26, 2024 @ 11:59pm

Please read the instructions carefully before submitting your assignment.

1. This assignment requires you to:
  - Upload your Quarto markdown files to a **git** repository
  - Upload a **PDF** file on Canvas
2. Don't collapse any code cells before submitting.
3. Remember to make sure all your code output is rendered properly before uploading your submission.


Please add your name to the the author information in the frontmatter before submitting your assignment.

## Question 1

💡 20 points

In this question, we will walk through the process of *forking* a `git` repository and submitting a *pull request*.

1. Navigate to the Github repository [here](#) and fork it by clicking on the icon in the top right



Provide a sensible name for your forked repository when prompted.

2. Clone your Github repository on your local machine

```
$ git clone <<insert your repository url here>>
$ cd hw-1
```

3. In order to activate the R environment for the homework, make sure you have `renv` installed beforehand. To activate the `renv` environment for this assignment, open an instance of the R console from within the directory and type

```
renv::activate()
```

Follow the instructions in order to make sure that `renv` is configured correctly.

4. Work on the *remaining part* of this assignment as a `.qmd` file.
  - Create a PDF and HTML file for your output by modifying the YAML frontmatter for the Quarto `.qmd` document
5. When you're done working on your assignment, push the changes to your github repository.
6. Navigate to the original Github repository [here](#) and submit a pull request linking to your repository.

Remember to **include your name** in the pull request information!

If you're stuck at any step along the way, you can refer to the [official Github docs here](#)

## Question 2

💡 30 points

Consider the following vector

```
my_vec <- c(
  "+0.07",
  "-0.07",
  "+0.25",
  "-0.84",
  "+0.32",
  "-0.24",
  "-0.97",
  "-0.36",
  "+1.76",
  "-0.36"
)
```

For the following questions, provide your answers in a code cell.

1. What data type does the vector contain?

```
# It contains character data
```

1. Create two new vectors called `my_vec_double` and `my_vec_int` which converts `my_vec` to Double & Integer types, respectively,

```
my_vec_double <- as.double(my_vec)

my_vec_int <- as.integer(my_vec)
```

1. Create a new vector `my_vec_bool` which comprises of:
  - TRUE if an element in `my_vec_double` is  $\leq 0$
  - FALSE if an element in `my_vec_double` is  $\geq 0$

```
my_vec_bool <- c()
for (x in 1:length(my_vec_double)) {
  if (my_vec_double[x]>0) {
    my_vec_bool[x] <- T
  }
  else {
    my_vec_bool[x] <- F
  }
}
```


How many elements of `my\_vec\_double` are greater than zero?

# 4

1. Sort the values of `my_vec_double` in ascending order.

```
correct_order <- order(my_vec_double)
my_vec_double <- my_vec_double[correct_order]
```

### Question 3

 50 points

In this question we will get a better understanding of how R handles large data structures in memory.

1. Provide R code to construct the following matrices:

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix} \quad \text{and} \quad \begin{bmatrix} 1 & 2 & 3 & 4 & 5 & \dots & 100 \\ 1 & 4 & 9 & 16 & 25 & \dots & 10000 \end{bmatrix}$$

## Warning

```
A <- matrix(  
  c(1,2,3,4,5,6,7,8,9),  
  nrow=3,  
  ncol=3,  
  byrow=T  
)
```

A

	[,1]	[,2]	[,3]
[1,]	1	2	3
[2,]	4	5	6
[3,]	7	8	9

```
row1 <- c(1:100)  
row2 <- row1^2  
B <- matrix(  
  c(row1,row2),  
  nrow=2,  
  ncol=length(row1),  
  byrow=T  
)
```

B

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]	[,11]	[,12]	[,13]	[,14]
[1,]	1	2	3	4	5	6	7	8	9	10	11	12	13	14
[2,]	1	4	9	16	25	36	49	64	81	100	121	144	169	196
	[,15]	[,16]	[,17]	[,18]	[,19]	[,20]	[,21]	[,22]	[,23]	[,24]	[,25]	[,26]		
[1,]	15	16	17	18	19	20	21	22	23	24	25	26		
[2,]	225	256	289	324	361	400	441	484	529	576	625	676		
	[,27]	[,28]	[,29]	[,30]	[,31]	[,32]	[,33]	[,34]	[,35]	[,36]	[,37]	[,38]		
[1,]	27	28	29	30	31	32	33	34	35	36	37	38		
[2,]	729	784	841	900	961	1024	1089	1156	1225	1296	1369	1444		
	[,39]	[,40]	[,41]	[,42]	[,43]	[,44]	[,45]	[,46]	[,47]	[,48]	[,49]	[,50]		
[1,]	39	40	41	42	43	44	45	46	47	48	49	50		
[2,]	1521	1600	1681	1764	1849	1936	2025	2116	2209	2304	2401	2500		
	[,51]	[,52]	[,53]	[,54]	[,55]	[,56]	[,57]	[,58]	[,59]	[,60]	[,61]	[,62]		
[1,]	51	52	53	54	55	56	57	58	59	60	61	62		
[2,]	2601	2704	2809	2916	3025	3136	3249	3364	3481	3600	3721	3844		

```

      [,63] [,64] [,65] [,66] [,67] [,68] [,69] [,70] [,71] [,72] [,73] [,74]
[1,]    63    64    65    66    67    68    69    70    71    72    73    74
[2,]  3969  4096  4225  4356  4489  4624  4761  4900  5041  5184  5329  5476
      [,75] [,76] [,77] [,78] [,79] [,80] [,81] [,82] [,83] [,84] [,85] [,86]
[1,]    75    76    77    78    79    80    81    82    83    84    85    86
[2,]  5625  5776  5929  6084  6241  6400  6561  6724  6889  7056  7225  7396
      [,87] [,88] [,89] [,90] [,91] [,92] [,93] [,94] [,95] [,96] [,97] [,98]
[1,]    87    88    89    90    91    92    93    94    95    96    97    98
[2,]  7569  7744  7921  8100  8281  8464  8649  8836  9025  9216  9409  9604
      [,99] [,100]
[1,]    99    100
[2,]  9801 10000

```

### Tip

Recall the discussion in class on how R fills in matrices

In the next part, we will discover how knowledge of the way in which a matrix is stored in memory can inform better code choices. To this end, the following function takes an input  $n$  and creates an  $n \times n$  matrix with random entries.

```

generate_matrix <- function(n){
  return(
    matrix(
      rnorm(n^2),
      nrow=n
    )
  )
}

```

For example:

```
generate_matrix(4)
```

```

      [,1]      [,2]      [,3]      [,4]
[1,]  2.0542648 0.08078202 -1.6265563  1.9874738
[2,] -1.7403692 0.02103011  0.9440185  0.2071361
[3,] -0.2449517 0.55199875  0.6260041  1.4016325
[4,]  0.8953322 0.16265631 -2.1974268  0.5522514

```

Let  $M$  be a fixed  $50 \times 50$  matrix

```
M <- generate_matrix(5000)
mean(M)
```

```
[1] 0.0002103701
```

2. Write a function `row_wise_scan` which scans the entries of `M` one row after another and outputs the number of elements whose value is  $\geq 0$ . You can use the following **starter code**

```
row_wise_scan <- function(x){
  n <- nrow(x)
  m <- ncol(x)

  # Insert your code here

  count <- 0
  for(i in 1:n){
    for(j in 1:m){
      if(x[i,j]>0){
        count <- count + 1
      }
    }
  }

  return(count)
}

row_wise_scan(M)
```

```
[1] 12502182
```

3. Similarly, write a function `col_wise_scan` which does exactly the same thing but scans the entries of `M` one column after another

```
col_wise_scan <- function(x){
  count <- 0
  n <- nrow(x)
  m <- ncol(x)

  # Insert your code here
```

```

    for(j in 1:n){
      for(i in 1:m){
        if(x[i,j]>0){
          count <- count + 1
        }
      }
    }

    return(count)
  }

col_wise_scan(M)

```

```
[1] 12502182
```

You can check if your code is doing what it's supposed to using the function here<sup>1</sup>

4. Between `col_wise_scan` and `row_wise_scan`, which function do you expect to take shorter to run? Why? I expect `col_wise_scan` to be shorter because the function is iterating through a single vector rather than switching vectors every time.

---

<sup>1</sup>If your code is right, the following code should evaluate to be TRUE

```

library(tidyverse)

-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr      1.1.4      v readr      2.1.5
v forcats    1.0.0      v stringr    1.5.1
v ggplot2    3.4.4      v tibble     3.2.1
v lubridate  1.9.3      v tidyr      1.3.0
v purrr      1.0.2
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()     masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

sapply(1:100, function(i) {
  x <- generate_matrix(100)
  row_wise_scan(x) == col_wise_scan(x)
}) %>% sum == 100

```

```
[1] TRUE
```



5. Write a function `time_scan` which takes in a method `f` and a matrix `M` and outputs the amount of time taken to run `f(M)`

```
time_scan <- function(f, m){  
  initial_time <- Sys.time() # Write your code here  
  f(m)  
  final_time <- Sys.time() # Write your code here  
  
  total_time_taken <- final_time - initial_time  
  return(total_time_taken)  
}
```

Provide your output to

```
list(  
  row_wise_time = time_scan(row_wise_scan, M),  
  col_wise_time = time_scan(col_wise_scan, M)  
)
```

```
$row_wise_time  
Time difference of 1.741782 secs
```

```
$col_wise_time  
Time difference of 1.952778 secs
```

Which took longer to run? ::: {.cell}

```
# row_wise_scan()
```

:::

6. Repeat this experiment now when:

- `M` is a  $100 \times 100$  matrix
- `M` is a  $1000 \times 1000$  matrix
- `M` is a  $5000 \times 5000$  matrix

What can you conclude? I can conclude that a column wise scan is going to be faster than a row wise scan.

## Appendix

Print your R session information using the following command

```
sessionInfo()
```

```
R version 4.2.2 (2022-10-31 ucrt)
Platform: x86_64-w64-mingw32/x64 (64-bit)
Running under: Windows 10 x64 (build 22621)
```

```
Matrix products: default
```

```
locale:
```

```
[1] LC_COLLATE=English_United States.utf8
[2] LC_CTYPE=English_United States.utf8
[3] LC_MONETARY=English_United States.utf8
[4] LC_NUMERIC=C
[5] LC_TIME=English_United States.utf8
```

```
attached base packages:
```

```
[1] stats      graphics  grDevices datasets  utils      methods    base
```

```
loaded via a namespace (and not attached):
```

```
[1] compiler_4.2.2    fastmap_1.1.1      cli_3.6.2           htmltools_0.5.7
[5] tools_4.2.2       rstudioapi_0.15.0  yaml_2.3.8          rmarkdown_2.25
[9] knitr_1.45        jsonlite_1.8.8     xfun_0.41           digest_0.6.34
[13] rlang_1.1.3       renv_1.0.3         evaluate_0.23
```