

Homework 1

Tisya Vaidya

Table of contents

Question 1	2
Question 2	3
Question 3	5

Appendix	10
-----------------	-----------

[Link to the Github repository](#)

! Due: Fri, Jan 26, 2024 @ 11:59pm

Please read the instructions carefully before submitting your assignment.

1. This assignment requires you to:
 - Upload your Quarto markdown files to a `git` repository
 - Upload a PDF file on Canvas
2. Don't collapse any code cells before submitting.
3. Remember to make sure all your code output is rendered properly before uploading your submission.

Please add your name to the the author information in the frontmatter before submitting your assignment.

Question 1

💡 20 points

In this question, we will walk through the process of *forking* a `git` repository and submitting a *pull request*.

1. Navigate to the Github repository [here](#) and fork it by clicking on the icon in the top right



Provide a sensible name for your forked repository when prompted.

2. Clone your Github repository on your local machine

```
$ git clone <<insert your repository url here>>
$ cd hw-1
```

3. In order to activate the R environment for the homework, make sure you have `renv` installed beforehand. To activate the `renv` environment for this assignment, open an instance of the R console from within the directory and type

```
#install.packages("tidyverse")
#install.packages("yaml")
#install.packages("renv")
renv::activate()
```

Follow the instructions in order to make sure that ``renv`` is configured correctly.


4. Work on the *remaining part* of this assignment as a `.qmd` file.
 - Create a PDF and HTML file for your output by modifying the YAML frontmatter for the Quarto `.qmd` document
5. When you're done working on your assignment, push the changes to your github repository.

6. Navigate to the original Github repository [here](#) and submit a pull request linking to your repository.

Remember to **include your name** in the pull request information!

If you're stuck at any step along the way, you can refer to the [official Github docs here](#)

Question 2

 30 points

Consider the following vector

```
my_vec <- c(
  "+0.07",
  "-0.07",
  "+0.25",
  "-0.84",
  "+0.32",
  "-0.24",
  "-0.97",
  "-0.36",
  "+1.76",
  "-0.36"
)
```

For the following questions, provide your answers in a code cell.

1. What data type does the vector contain? ::: {.cell}

```
typeof(my_vec)
```

```
[1] "character"
```

:::

1. Create two new vectors called `my_vec_double` and `my_vec_int` which converts `my_vec` to Double & Integer types, respectively,

```
# Create a new vector of type double
my_vec_double <- as.numeric(my_vec)
my_vec_double
```

```
[1] 0.07 -0.07 0.25 -0.84 0.32 -0.24 -0.97 -0.36 1.76 -0.36
```

```
# Create a new vector of type integer
my_vec_int <- as.integer(my_vec)
my_vec_int
```

```
[1] 0 0 0 0 0 0 0 0 1 0
```

1. Create a new vector `my_vec_bool` which comprises of:

- TRUE if an element in `my_vec_double` is ≤ 0
- FALSE if an element in `my_vec_double` is ≥ 0

How many elements of `my_vec_double` are greater than zero? ::: {.cell}

```
# Convert to double type first
my_vec_double <- as.numeric(my_vec)

# Create a logical vector
my_vec_bool <- my_vec_double <= 0

# Count elements greater than zero
sum(my_vec_double > 0)
```

```
[1] 4
```

::: 1. Sort the values of `my_vec_double` in ascending order.

```
# Sort the vector
sorted_vec <- sort(my_vec_double)
sorted_vec
```

```
[1] -0.97 -0.84 -0.36 -0.36 -0.24 -0.07 0.07 0.25 0.32 1.76
```

Question 3

💡 50 points

In this question we will get a better understanding of how R handles large data structures in memory.

1. Provide R code to construct the following matrices:

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix} \quad \text{and} \quad \begin{bmatrix} 1 & 2 & 3 & 4 & 5 & \dots & 100 \\ 1 & 4 & 9 & 16 & 25 & \dots & 10000 \end{bmatrix}$$

::: {.callout-warning} ## Tip

Recall the discussion in class on how R fills in matrices :::

In the next part, we will discover how knowledge of the way in which a matrix is stored in memory can inform better code choices. To this end, the following function takes an input n and creates an $n \times n$ matrix with random entries.

```
generate_matrix <- function(n){  
  return(  
    matrix(  
      rnorm(n^2),  
      nrow=n  
    )  
  )  
}
```

For example:

```
generate_matrix(4)
```

```
      [,1]      [,2]      [,3]      [,4]  
[1,] -0.5856412 -0.5588938 -0.49455369  0.07863609  
[2,] -1.1399401 -0.9997287 -1.52764179  0.54556752  
[3,] -0.5612194 -0.1763169 -0.09906992 -0.79476193  
[4,]  0.8720956  0.3131394 -1.12536155  0.58942771
```

Let M be a fixed 50×50 matrix

```
M <- generate_matrix(50)
mean(M)
```

```
[1] -0.01579699
```

```
# Matrix 1: 3x3 Matrix
matrix_1 <- matrix(1:9, nrow = 3, byrow = TRUE)

# Generate Matrix 2 with the specified pattern
matrix_2 <- matrix(c(1:100, (1:100)^2), nrow = 2, byrow= TRUE)

# Display matrices
print(matrix_1)
```

```
      [,1] [,2] [,3]
[1,]     1     2     3
[2,]     4     5     6
[3,]     7     8     9
```

```
print(matrix_2)
```

```
      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13] [,14]
[1,]     1     2     3     4     5     6     7     8     9     10     11     12     13     14
[2,]     1     4     9    16    25    36    49    64    81    100    121    144    169    196
      [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22] [,23] [,24] [,25] [,26]
[1,]    15    16    17    18    19    20    21    22    23    24    25    26
[2,]   225   256   289   324   361   400   441   484   529   576   625   676
      [,27] [,28] [,29] [,30] [,31] [,32] [,33] [,34] [,35] [,36] [,37] [,38]
[1,]    27    28    29    30    31    32    33    34    35    36    37    38
[2,]   729   784   841   900   961  1024  1089  1156  1225  1296  1369  1444
      [,39] [,40] [,41] [,42] [,43] [,44] [,45] [,46] [,47] [,48] [,49] [,50]
[1,]    39    40    41    42    43    44    45    46    47    48    49    50
[2,]  1521  1600  1681  1764  1849  1936  2025  2116  2209  2304  2401  2500
      [,51] [,52] [,53] [,54] [,55] [,56] [,57] [,58] [,59] [,60] [,61] [,62]
[1,]    51    52    53    54    55    56    57    58    59    60    61    62
[2,]  2601  2704  2809  2916  3025  3136  3249  3364  3481  3600  3721  3844
      [,63] [,64] [,65] [,66] [,67] [,68] [,69] [,70] [,71] [,72] [,73] [,74]
[1,]    63    64    65    66    67    68    69    70    71    72    73    74
```

```

[2,] 3969 4096 4225 4356 4489 4624 4761 4900 5041 5184 5329 5476
     [,75] [,76] [,77] [,78] [,79] [,80] [,81] [,82] [,83] [,84] [,85] [,86]
[1,] 75    76    77    78    79    80    81    82    83    84    85    86
[2,] 5625 5776 5929 6084 6241 6400 6561 6724 6889 7056 7225 7396
     [,87] [,88] [,89] [,90] [,91] [,92] [,93] [,94] [,95] [,96] [,97] [,98]
[1,] 87    88    89    90    91    92    93    94    95    96    97    98
[2,] 7569 7744 7921 8100 8281 8464 8649 8836 9025 9216 9409 9604
     [,99] [,100]
[1,] 99    100
[2,] 9801 10000

```

2. Write a function `row_wise_scan` which scans the entries of `M` one row after another and outputs the number of elements whose value is ≥ 0 . You can use the following **starter code**

```

row_wise_scan <- function(x){
  n <- nrow(x)
  m <- ncol(x)

  # Insert your code here
  count <- 0
  for(i in 1:n){
    for(j in 1:m){
      if(x[i, j] >= 0){
        count <- count + 1
      }
    }
  }

  return(count)
}

```

3. Similarly, write a function `col_wise_scan` which does exactly the same thing but scans the entries of `M` one column after another

```

col_wise_scan <- function(x){
  count <- 0
  # Insert your code here
  n <- nrow(x)
  m <- ncol(x)
  for(j in 1:m){
    for(i in 1:n){

```

```

        if(x[i, j] >= 0){
            count <- count + 1
        }
    }
}

return(count)
}

```

You can check if your code is doing what it's supposed to using the function here¹

4. Between `col_wise_scan` and `row_wise_scan`, which function do you expect to take shorter to run? Why?

Expectation on Runtime:

The `row_wise_scan` function is expected to run faster than `col_wise_scan`. This is because R stores matrices in a column-major order, so accessing elements column-wise results in better memory locality.

5. Write a function `time_scan` which takes in a method `f` and a matrix `M` and outputs the amount of time taken to run `f(M)`

```

time_scan <- function(f, M){
    initial_time <- Sys.time() # Write your code here
    f(M)
    final_time <- Sys.time() # Write your code here

    total_time_taken <- final_time - initial_time
    return(total_time_taken)
}

```

Provide your output to

```

list(
    row_wise_time = time_scan(row_wise_scan, M),
    col_wise_time = time_scan(col_wise_scan, M)
)

```

`$row_wise_time`

¹If your code is right, the following code should evaluate to be TRUE

```

{R sapply(1:100, function(i) {      x <- generate_matrix(100)      row_wise_scan(x) ==
col_wise_scan(x) }) %>% sum == 100

```


Time difference of 0.002946138 secs

\$col_wise_time

Time difference of 0.0001261234 secs

Which took longer to run? row_wise function took longer to run

6. Repeat this experiment now when:

- M is a 100×100 matrix
- M is a 1000×1000 matrix
- M is a 5000×5000 matrix

What can you conclude?

```
# Function to repeat experiment
repeat_experiment <- function(n) {
  M <- generate_matrix(n)
  time_results <- list(
    row_wise_time = time_scan(row_wise_scan, M),
    col_wise_time = time_scan(col_wise_scan, M)
  )
  return(time_results)
}

# Experiment for different matrix sizes
experiment_results <- list(
  matrix_100 = repeat_experiment(100),
  matrix_1000 = repeat_experiment(1000),
  matrix_5000 = repeat_experiment(5000)
)
experiment_results
```

\$matrix_100

\$matrix_100\$row_wise_time

Time difference of 0.0004849434 secs

\$matrix_100\$col_wise_time

Time difference of 0.003509045 secs

\$matrix_1000

\$matrix_1000\$row_wise_time

Time difference of 0.04827285 secs

\$matrix_1000\$col_wise_time

Time difference of 0.04623485 secs

\$matrix_5000

\$matrix_5000\$row_wise_time

Time difference of 1.526564 secs

\$matrix_5000\$col_wise_time

Time difference of 1.15332 secs

We can conclude that row_wise function takes longer to run than col_wise

Appendix

Print your R session information using the following command

```
sessionInfo()
```

R version 4.3.1 (2023-06-16)

Platform: aarch64-apple-darwin20 (64-bit)

Running under: macOS Monterey 12.0.1

Matrix products: default

BLAS: /Library/Frameworks/R.framework/Versions/4.3-arm64/Resources/lib/libRblas.0.dylib

LAPACK: /Library/Frameworks/R.framework/Versions/4.3-arm64/Resources/lib/libRlapack.dylib;

locale:

[1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8

time zone: America/New_York

tzcode source: internal

attached base packages:

[1] stats graphics grDevices datasets utils methods base

loaded via a namespace (and not attached):

[1] compiler_4.3.1	fastmap_1.1.1	cli_3.6.2	htmltools_0.5.7
[5] tools_4.3.1	rstudioapi_0.15.0	yaml_2.3.8	rmarkdown_2.25
[9] knitr_1.45	jsonlite_1.8.8	xfun_0.41	digest_0.6.34
[13] rlang_1.1.3	renv_1.0.3	evaluate_0.23	